

Inference-side efficiency: SR vs. Token cost

